
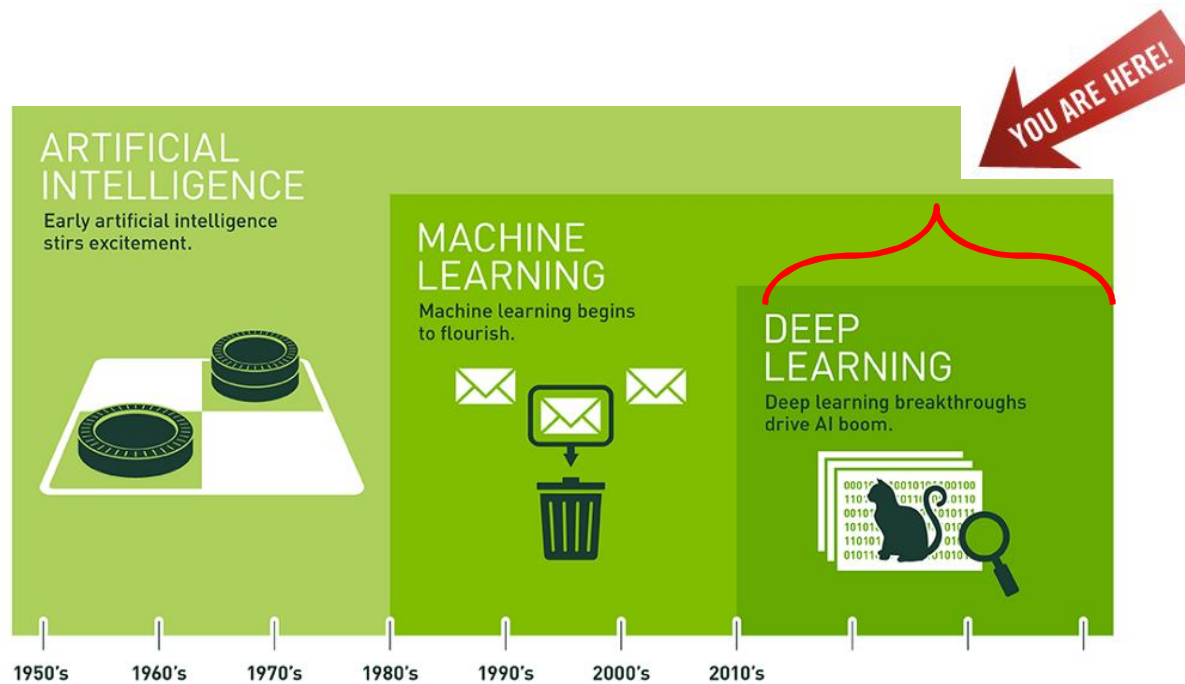

Artificial Intelligence: Deep Learning

(too recent to be in the Russell & Norvig book)
many slides from: Y. Bengio, A. Ng and Y. LeCun

Today

1. Motivation 
2. Feature Learning
3. CNNs for Image Processing
4. Conclusion

History of AI




Major Breakthroughs

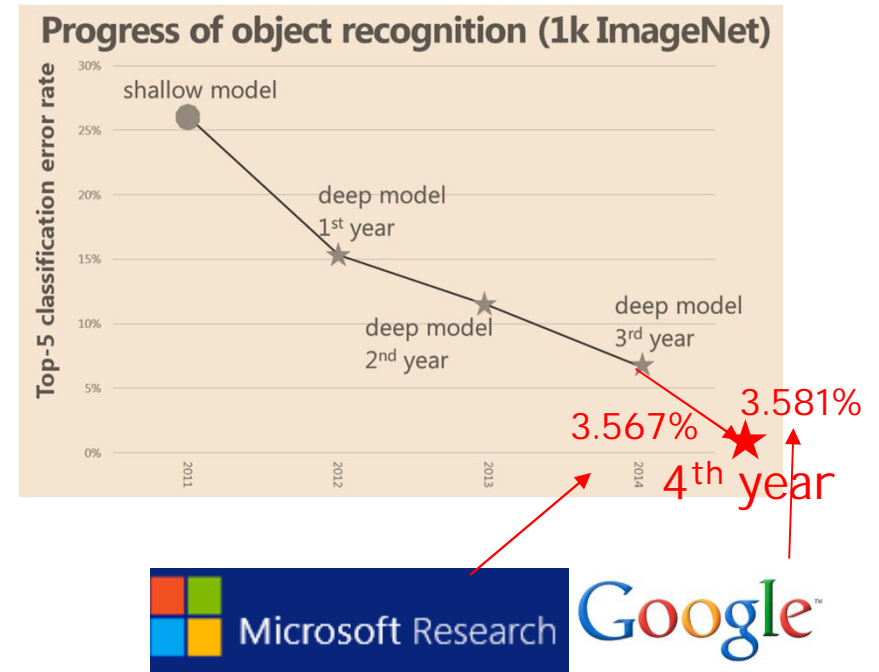
- Speech Recognition & Machine Translation (2010+)
- Image Recognition & Computer Vision (2012+)



Object recognition



Self driving cars



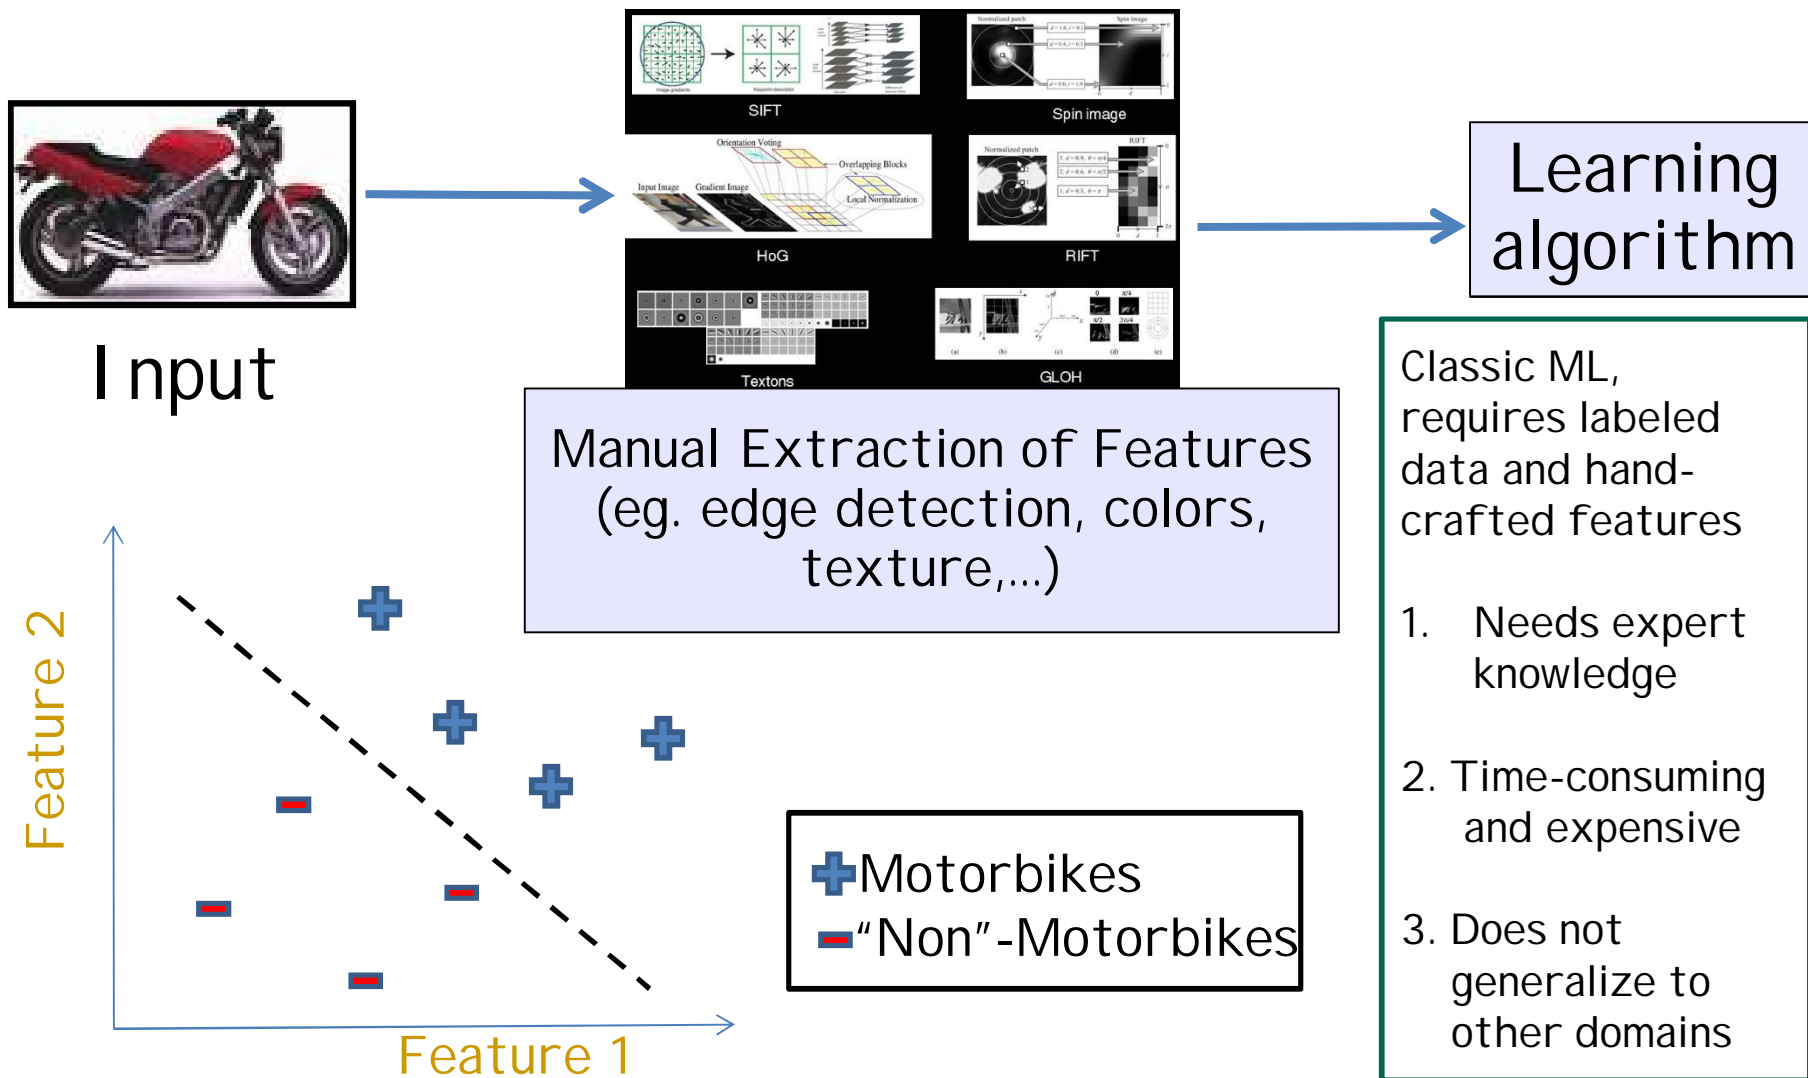
- Natural Language Processing (2014+)
- ...

Today

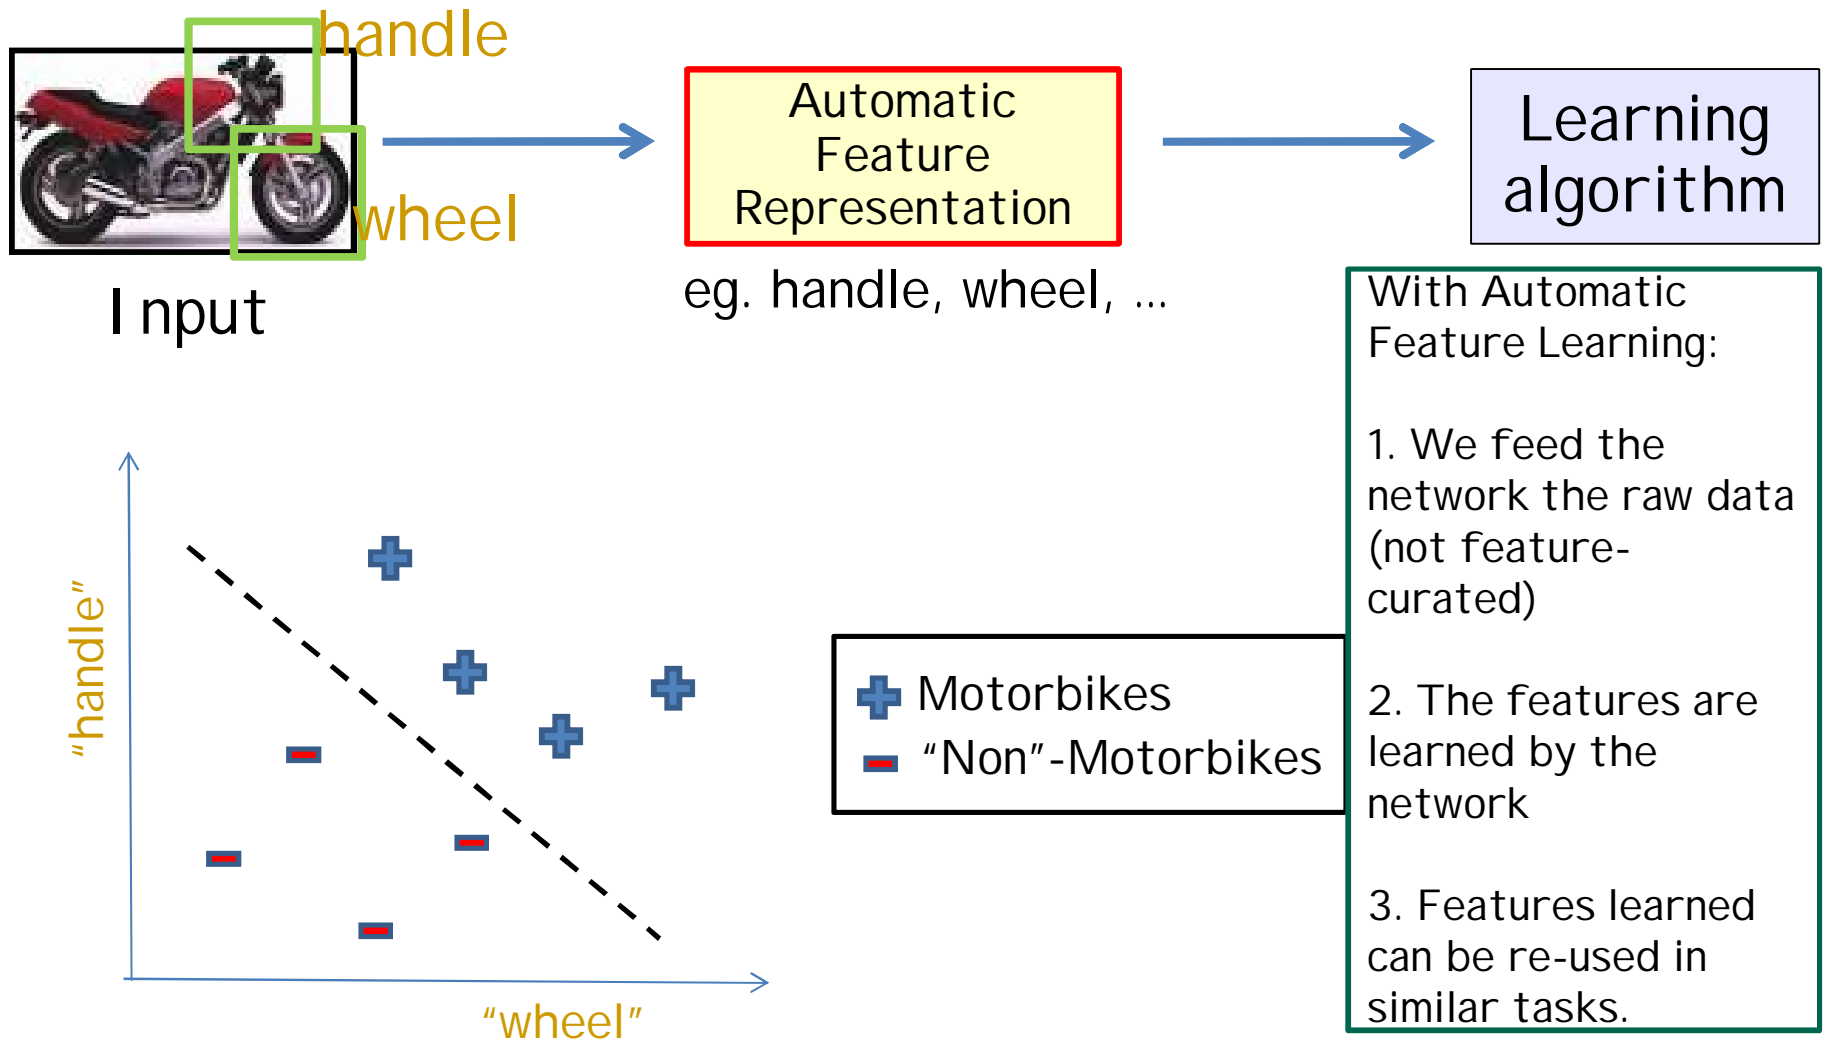
1. Motivation
2. Feature Learning
3. CNNs for Image Processing
4. Conclusion




Classic ML



Automatic Feature Learning



Advantages of Unsupervised Feature Learning



SPEECH LANGUAGE THERAPY IN THE EDUCATIONAL CONTEXT
 BSc Student: Alexandra Paula Rodrigues
 Course: English Language
 University of the South Atlantic - Universidade do Sul da América
 Email: alexandra.paula@unisa.edu.br

INTRODUCTION



Abstract: This paper aims to explore the role of speech language therapy in the educational context, focusing on the challenges faced by students with communication disorders and the strategies used by professionals to address these challenges.

Keywords: Speech language therapy, Educational context, Communication disorders, Students, Professionals, Strategies.

Introduction: The purpose of this paper is to explore the role of speech language therapy in the educational context, focusing on the challenges faced by students with communication disorders and the strategies used by professionals to address these challenges.

Conclusion: The role of speech language therapy in the educational context is crucial for addressing the communication needs of students with disorders and promoting their academic and social development.

References: American Speech-Language-Hearing Association. (2015). *Principles and Practices of Speech-Language Pathology*. Silver Spring, MD: ASHA.



Automatic Feature Learning

Deep Learning = Machine learning algorithms based on learning multiple levels of representation / abstraction.

– Y. Bengio


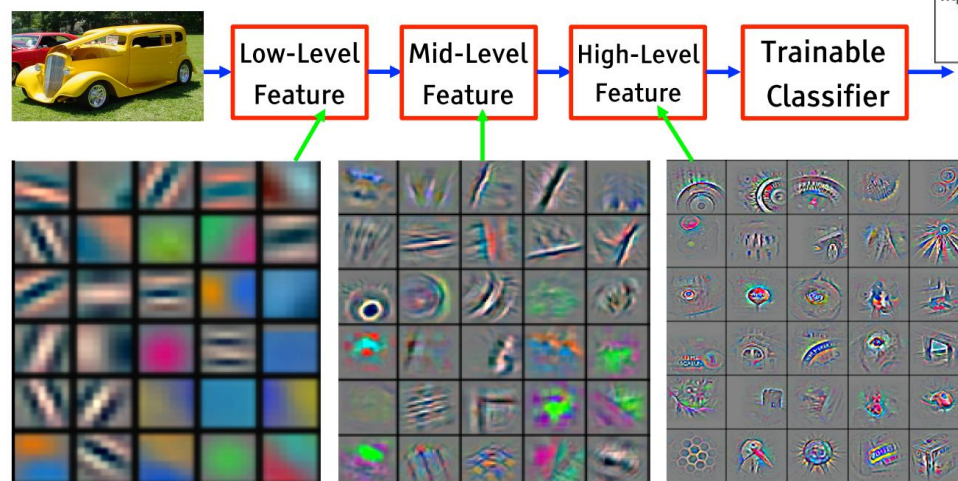
- ❑ Each layer learns more abstract features that are then combined / composed into higher-level features automatically
- ❑ Like the human brain ...
 - ❑ has many layers of neurons which act as feature detectors
 - ❑ detecting more and more abstract features as you go up
- ❑ E.g. to classify an image of a cat:
 - Bottom Layers: Edge detectors, curves, corners straight lines
 - Middle Layers: Fur patterns, eyes, ears
 - Higher Layers: Body, head, legs
 - Top Layer: Cat or Dog



What Types of Features?

- For image recognition

- pixel -> edge -> texton -> motif -> part -> object



- For NLP

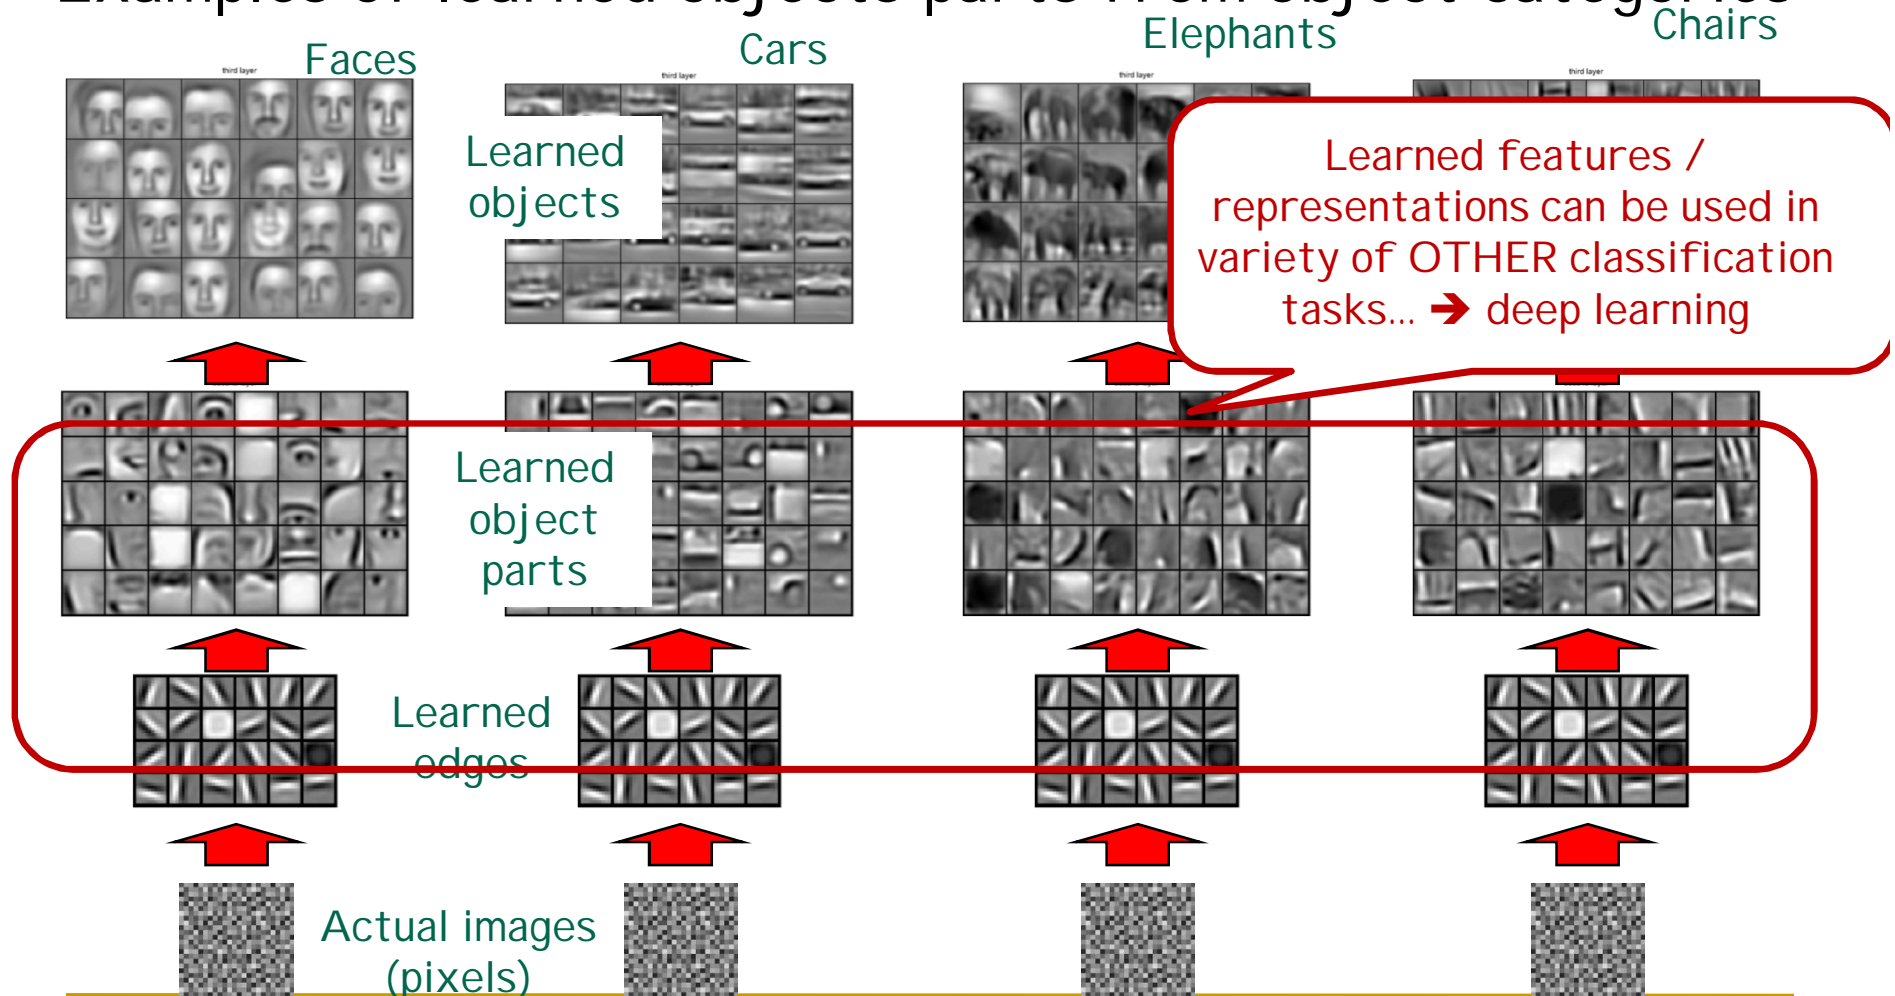
- character -> word -> constituents -> clause -> sentence -> discourse

- For speech:

- sample -> spectral band -> sound -> ... phone -> phoneme -> word

Eg: Learning Image Features

Examples of learned objects parts from object categories

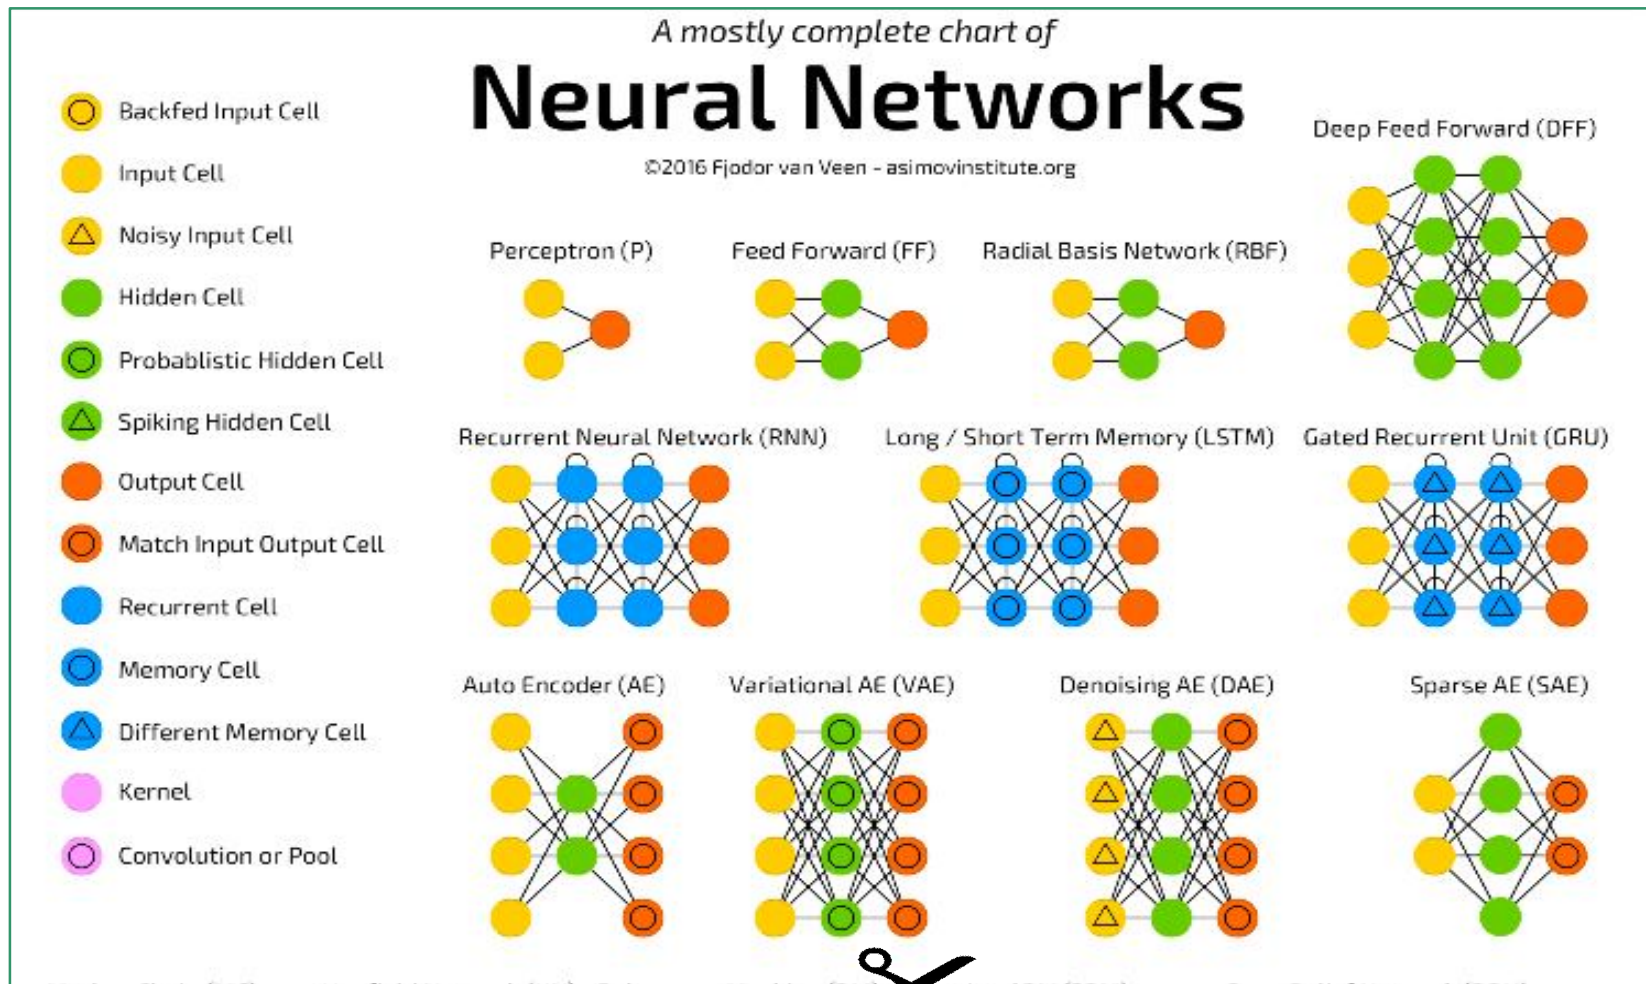


Today

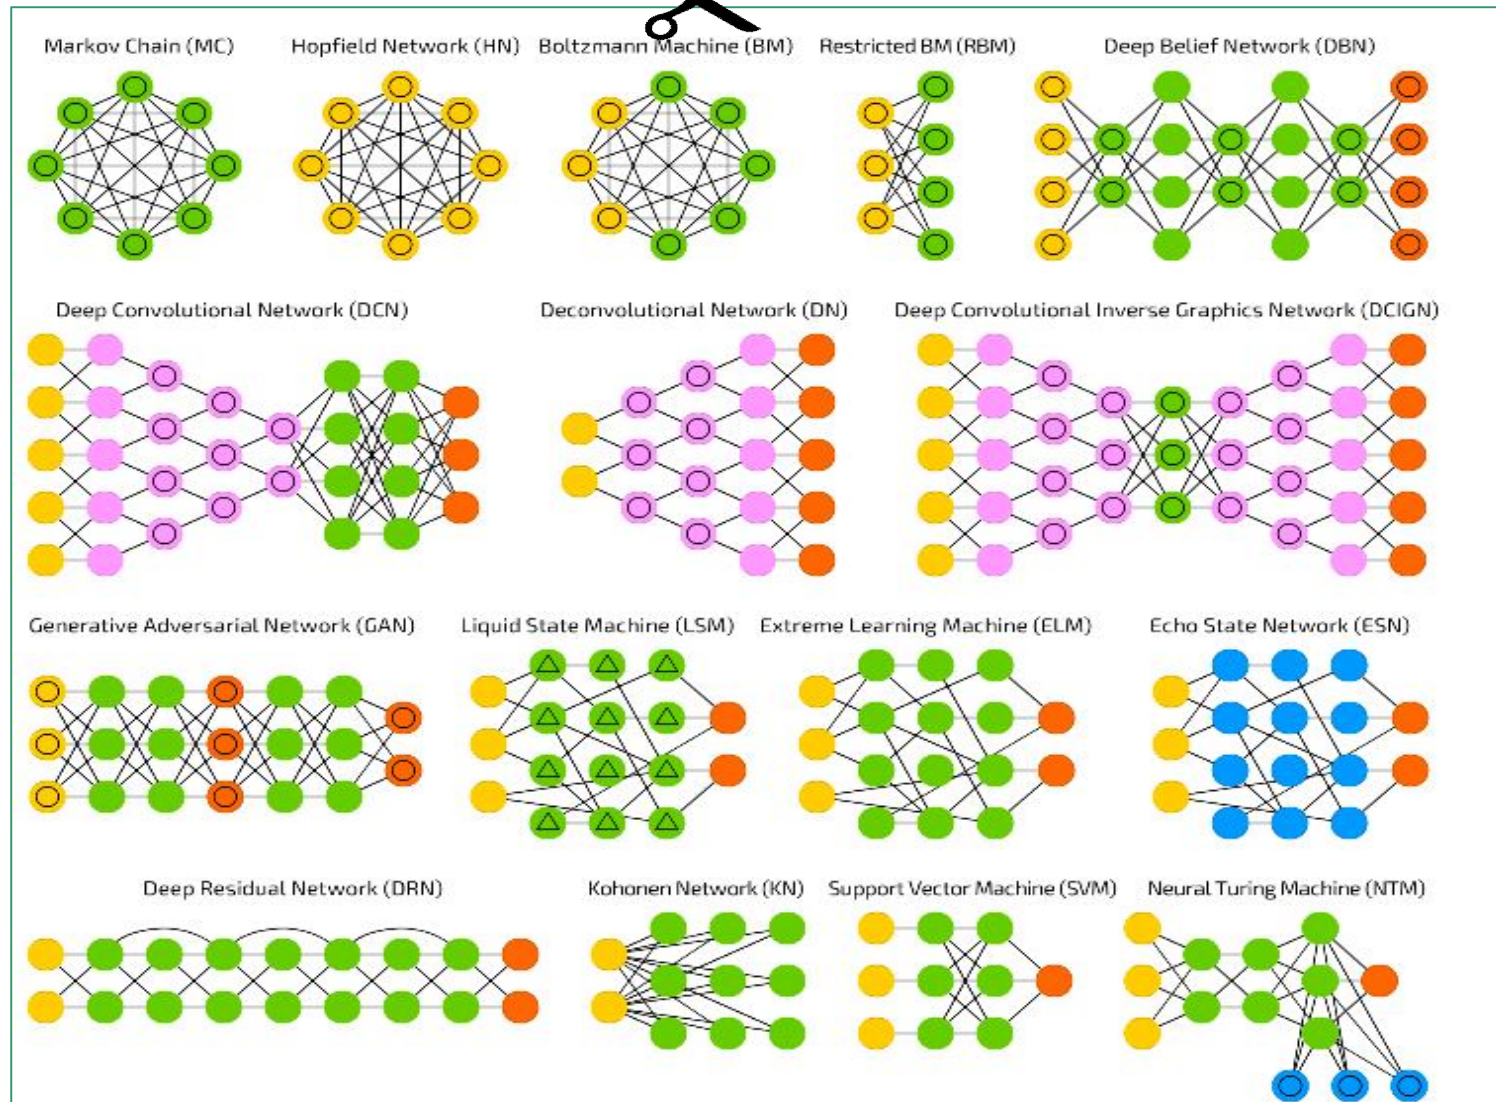
1. Motivation
2. Feature Learning
3. CNNs for Image Processing
4. Conclusion



Many Types of Neural Networks





Many Types of Deep Networks (con't)



CNNs for Image Processing

CNNs = Convolutional Neural Networks



| | | | |
|-----|-----|-----|-----|
| 25 | 2 | 1 | 44 |
| 223 | 7 | 6 | 60 |
| 196 | 8 | 2 | 148 |
| 249 | 1 | 3 | 40 |
| 60 | 7 | 1 | 154 |
| 59 | 1 | 7 | 213 |
| 214 | 7 | 3 | 163 |
| 89 | 182 | 219 | 13 |
| 74 | 146 | 113 | 72 |
| 89 | 18 | 244 | 85 |
| 1 | 4 | 8 | 97 |
| 3 | 4 | 2 | 121 |
| 2 | 1 | 2 | 131 |
| 7 | 6 | 8 | 47 |
| 3 | 5 | 5 | 126 |
| 7 | 6 | 8 | 121 |
| 5 | 3 | 1 | 237 |


Image of a 4 in grey scale
Value = 0-> white 255->black

CNNs for Image Processing

- Standard input of the image in the ANN:

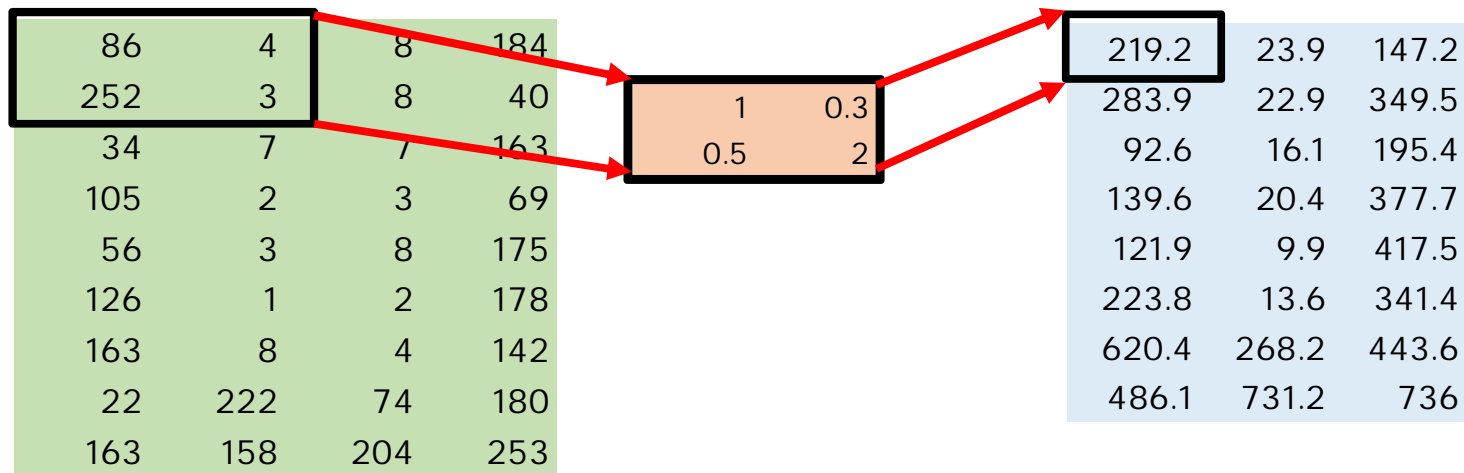
- 1 pixel = 1 input.
- Eg. color image of 200x200x 3channels (RGB)
 - > in a fully connected ANN, a neuron of the input layer has $200 \times 200 \times 3 = 120,000$ weights
 - > huge number of parameters, can easily overfit
- We linearize the image ==> We lose spatial information

| | | | |
|-----|-----|-----|-----|
| 25 | 2 | 1 | 44 |
| 223 | 7 | 6 | 60 |
| 196 | 8 | 2 | 148 |
| 249 | 1 | 3 | 40 |
| 60 | 7 | 1 | 154 |
| 59 | 1 | 7 | 213 |
| 214 | 7 | 3 | 163 |
| 89 | 182 | 219 | 13 |
| 74 | 146 | 113 | 72 |
| 89 | 18 | 244 | 85 |
| 1 | 4 | 8 | 97 |
| 3 | 4 | 2 | 121 |
| 2 | 1 | 2 | 131 |
| 7 | 6 | 8 | 47 |
| 3 | 5 | 5 | 126 |
| 7 | 6 | 8 | 121 |
| 5 | 3 | 1 | 237 |



Convolutional Layer

- Use a filter (aka kernel) that “convolves” on the image
- Filter = small weight matrix to learn



I (original image)


W (filter)

C (activation map)

$$\begin{aligned} C_{11} &= (I_{11} \times W_{11}) + (I_{12} \times W_{12}) + (I_{21} \times W_{21}) + (I_{22} \times W_{22}) \\ &= 86 \times 1 + 4 \times 0.3 + 252 \times 0.5 + 3 \times 2 = 219.2 \end{aligned}$$

Convolutional Layer

- Use a filter (aka kernel) that “convolves” on the image
- Filter = small weight matrix to learn



I (original image)


W (filter)

C (activation map)

$$\begin{aligned} C_{12} &= (I_{12} \times W_{11}) + (I_{13} \times W_{12}) + (I_{22} \times W_{21}) + (I_{23} \times W_{22}) \\ &= 252 \times 1 + 3 \times 0.3 + 34 \times 0.5 + 7 \times 2 = 283.9 \end{aligned}$$

Convolutional Layer

- Use a filter (aka kernel) that “convolves” on the image
- Filter = small weight matrix to learn



I (original image)

W (filter)

C (activation map)

$$\begin{aligned} C_{13} &= (I_{13} \times W_{11}) + (I_{14} \times W_{12}) + (I_{23} \times W_{21}) + (I_{24} \times W_{22}) \\ &= 34 \times 1 + 7 \times 0.3 + 105 \times 0.5 + 2 \times 2 = 92.6 \end{aligned}$$

Convolutional Layer

- Use a filter (aka kernel) that “convolves” on the image
- Filter = small weight matrix to learn

| | | | |
|-----|-----|-----|-----|
| 86 | 4 | 8 | 184 |
| 252 | 3 | 8 | 40 |
| 34 | 7 | 7 | 163 |
| 105 | 2 | 3 | 69 |
| 56 | 3 | 8 | 175 |
| 126 | 1 | 2 | 178 |
| 163 | 8 | 4 | 142 |
| 22 | 222 | 74 | 180 |
| 163 | 158 | 204 | 253 |

I (original image)

| | |
|-----|-----|
| 1 | 0.3 |
| 0.5 | 2 |

W (filter)

| | | |
|-------|-------|-------|
| 219.2 | 23.9 | 147.2 |
| 283.9 | 22.9 | 349.5 |
| 92.6 | 16.1 | 195.4 |
| 139.6 | 20.4 | 377.7 |
| 121.9 | 9.9 | 417.5 |
| 223.8 | 13.6 | 341.4 |
| 620.4 | 268.2 | 443.6 |
| 486.1 | 731.1 | 736 |

C (activation map)

$$\begin{aligned} C_{38} &= (I_{37} \times W_{11}) + (I_{47} \times W_{12}) + (I_{38} \times W_{21}) + (I_{37} \times W_{22}) \\ &= 74 \times 1 + 180 \times 0.3 + 204 \times 0.5 + 253 \times 2 = 736 \end{aligned}$$

Learn the Filters

| | | | | | |
|----|-----|-----|-----|-----|-----|
| 18 | 54 | 51 | 239 | 244 | 188 |
| 55 | 121 | 75 | 78 | 95 | 88 |
| 35 | 24 | 204 | 113 | 109 | 221 |
| 3 | 154 | 104 | 235 | 25 | 130 |
| 15 | 253 | 225 | 159 | 78 | 233 |
| 68 | 85 | 180 | 214 | 245 | 0 |

$I (6 \times 6)$

| | | |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

$W (3 \times 3)$

| | | | |
|-----|-----|-----|-----|
| 429 | 505 | 686 | 856 |
| 261 | 792 | 412 | 640 |
| 633 | 653 | 851 | 751 |
| 608 | 913 | 713 | 657 |


$C (4 \times 4)$

- The weight matrix (filter/kernel) behaves like a filter
- The network learns the values of the filter(s) that activate when they “see” some visual feature that is useful to identify the object (the final classification)
 - Ex. a horizontal line, a blotch of some color, a circle...

Convolution Hyper-parameters


1. Stride
2. Padding

Stride




I (7×7)

W (3×3) with stride = 1




C (5×5)



I (7×7)

W (3×3) with stride = 2



C (3×3)

Padding

| | | | |
|---|---|---|---|
| 9 | 0 | 0 | 1 |
| 1 | 0 | 1 | 0 |
| 0 | 1 | 1 | 2 |
| 2 | 1 | 0 | 1 |

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

| | |
|---|---|
| 0 | 1 |
| 1 | 1 |

9 not picked up ;-)

filter should pick up high values surrounded by low values

| | | | | | |
|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 9 | 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 1 | 2 | 0 |
| 0 | 2 | 1 | 0 | 1 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |

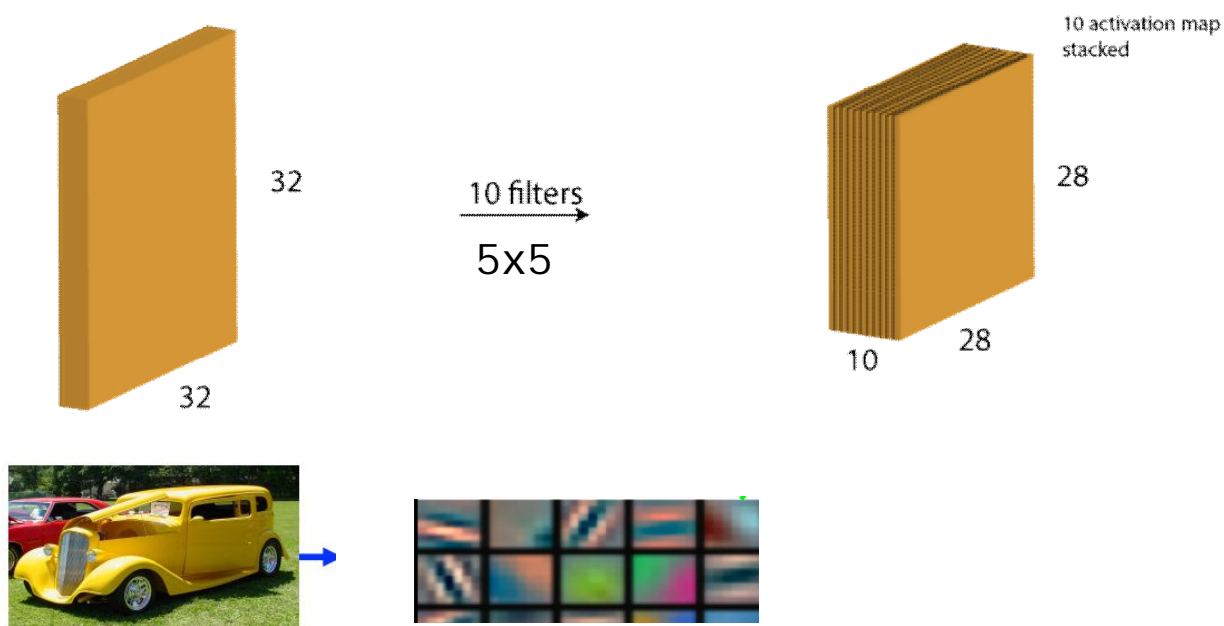
| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

| | | | |
|---|---|---|---|
| 9 | 0 | 0 | 1 |
| 1 | 0 | 1 | 0 |
| 0 | 1 | 1 | 2 |
| 2 | 1 | 0 | 1 |

9 picked up ;-)

Learn Several Filters

- So we create 1 activation map per filter



Pooling Layer

- Used to:
 - To reduce the size of the activation maps
 - So that we reduce the number of parameters of the network and hence avoid overfitting.

- Several strategies:

- Max pooling

| | | | | | | |
|-----|-----|-----|-----|---|-----|-----|
| 429 | 505 | 686 | 856 | → | 792 | 856 |
| 261 | 792 | 412 | 640 | | 913 | 851 |
| 633 | 653 | 851 | 751 | | | |
| 608 | 913 | 713 | 657 | | | |


- Average pooling

| | | | | | | |
|-----|-----|-----|-----|---|-------|-------|
| 429 | 505 | 686 | 856 | → | 496.8 | 648.5 |
| 261 | 792 | 412 | 640 | | 701.8 | 743 |
| 633 | 653 | 851 | 751 | | | |
| 608 | 913 | 713 | 657 | | | |

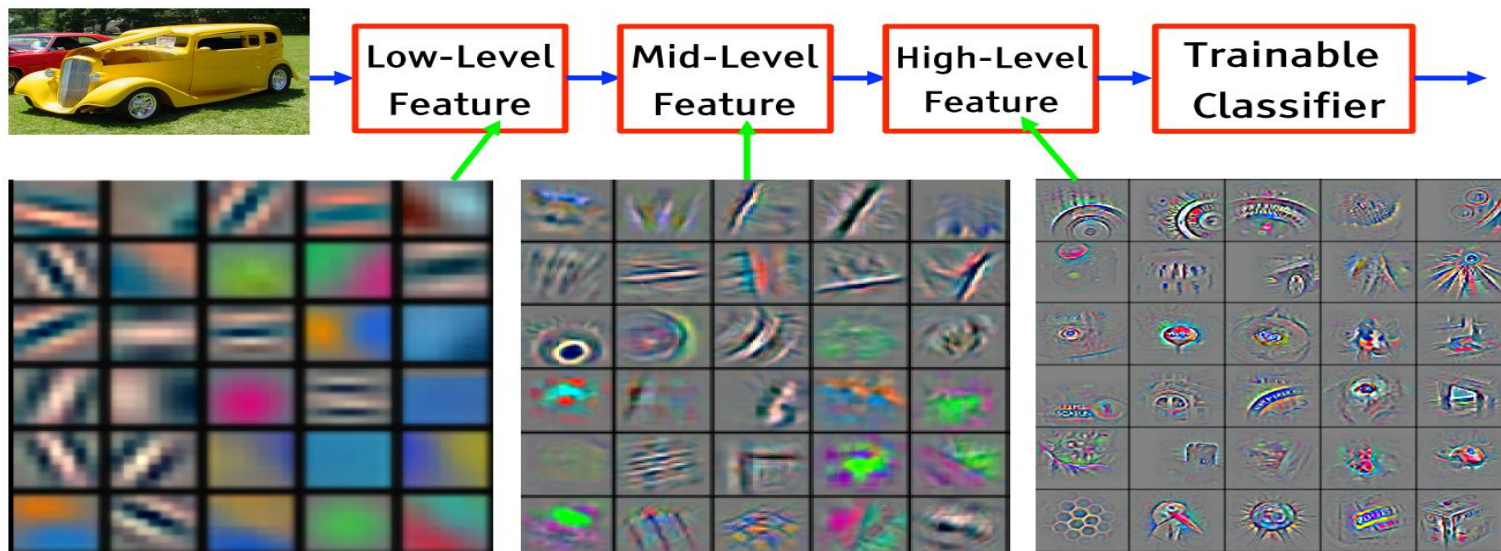

- ...

Architecture of a CNN


- Stack:
 - ▣ Convolutional Layers
 - ▣ Pooling Layers
- Finish off with:
 - ▣ A fully connected layer at the end for the final classification



Learning a Hierarchy of Features



Example of a CNN



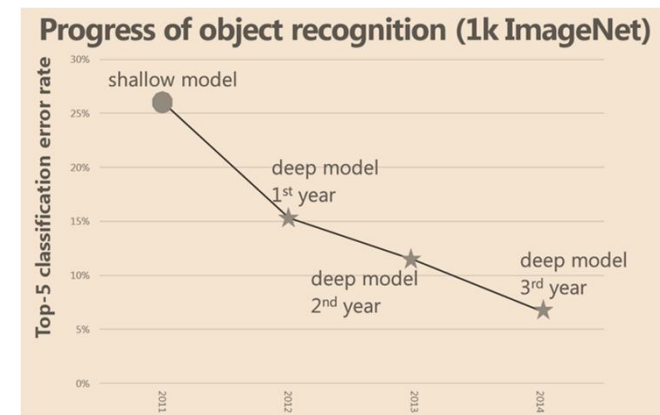
Successful CNN Networks

■ LeNet

- First successful applications of CNNs
- Developed by Yann LeCun in the 1990's
- used to read zip codes, digits, etc.

■ AlexNet

- First work that popularized CNNs for computer vision
- developed by Alex Krizhevsky, Ilya Sutskever and Geoff Hinton (U. of Toronto)
- In 2012 significantly outperformed all teams at the ImageNet ILSVRC challenge



Today

1. Motivation
2. Feature Learning
3. CNNs for Image Processing
4. Conclusion



Why now?

1. Basic science

- ❑ Backpropagation did not work / overfitting...
- ❑ *now*: developed method for training, better activation functions, better architectures....
- ❑ Need for lots training data...
- ❑ *now*: we have massive amounts + unsupervised pre-training

2. GPU computing

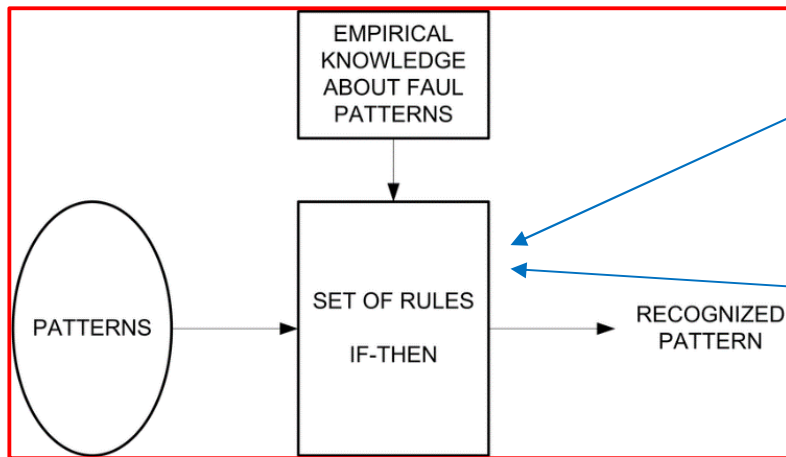
- ❑ Neural networks are very very long to train... (days, weeks)
- ❑ *now*: use of GPU's which are optimized for very fast matrix multiplication

3. Open Access to resources


- ❑ *now* : Access to DL methods, code and frameworks
- ❑ *now* : Fast turnaround from idea to implementation

History of AI

Rules written by experts ☹️
(eg. linguistics, medical doctors,...)



Rules hand-written by linguists

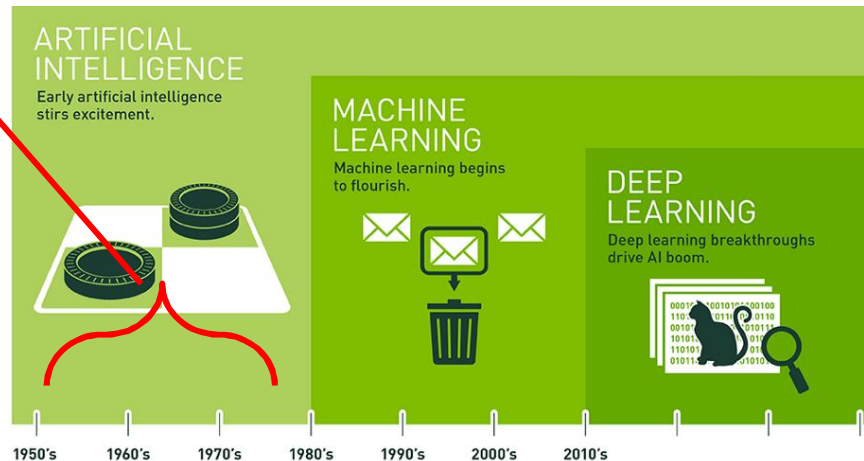


IF

- "yes" is equal to uniform_layer_flow
- THETA is greater than 45.0
- THETA is less than or equal to 90.0
- C4 is greater than $(Lm/(0.8*(Hs-H0)))$
- C6 is greater than $(Lb/(0.8*(Hs-H0)))$
- C9 is less than or equal to $(Lt/(0.8*(Hs-H0)))$

Then

- flow_type_ok is confirmed
- "V2" is assigned to flow_type
- "No" is assigned to wake_attachment
- Find coanda_attachment_value

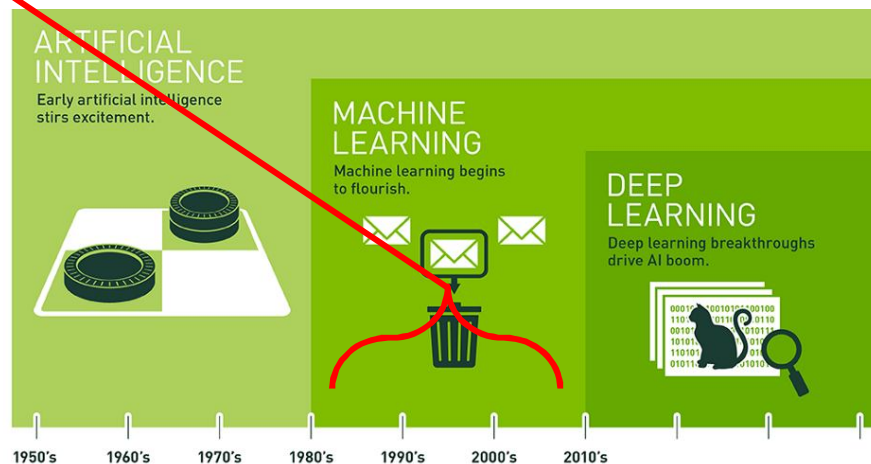
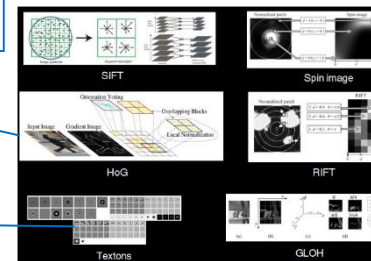




https://www.researchgate.net/profile/Dubravko_Miljkovic/publication/268239364/figure/fig30/AS:394719407427587@147111

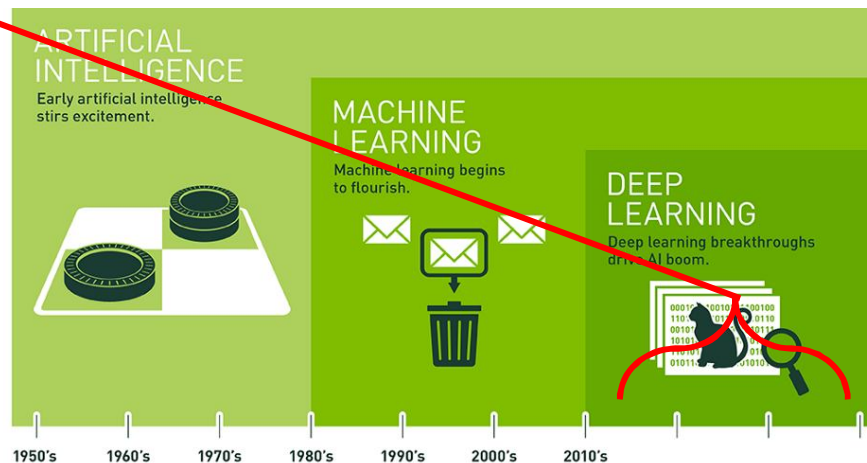

<http://www.cormix.info/images/RuleTreeExample.jpg>

History of AI

Rules learns via the data ;-) 😊
 But: features identified by the experts 😞
 (eg. linguistics, medical doctors,...)



History of AI



<https://www.linkedin.com/pulse/goedels-incompleteness-theorem-emergence-ai-eberhard-schoeneburg/>

Conclusion

- Deep Learning is thriving !
 - *vision*
 - *image processing*
 - *speech recognition*
 - *natural language processing*
 - ...
- Canada is a world leader in Deep Learning
 1. Montreal: (Bengio et al.) [MILA](#)
 2. Toronto: (Hinton et al.) [Vector Institute](#)
 3. Edmonton: (Sutton et al.) [AMLI](#)

